

1. Numerical error & complexity

Floating point: real numbers stored with finite precision; roundoff errors inevitable.

Machine epsilon: smallest $\varepsilon > 0$ such that $1 + \varepsilon \neq 1$ in floating point.

Floating point and error.

- Double precision: $\varepsilon_{\text{mach}} \approx 2.2 \times 10^{-16}$.
- Relative error: $e_{\text{rel}} = (\hat{x} - x)/x$; absolute error: $e_{\text{abs}} = \hat{x} - x$.
- Avoid subtracting nearly equal numbers (catastrophic cancellation).

Truncation error: error from approximating an infinite process (e.g. Taylor series, step size).

Big-O: $f(n) = O(g(n))$ if $|f(n)| \leq C|g(n)|$ for large n .

Big-O complexity.

- Dot product: $O(n)$.
- Matrix–vector: $O(mn)$.
- Matrix–matrix: $O(mnp)$.
- Dense linear solve: $O(n^3)$.

Example (catastrophic cancellation):

Compute $1 - \cos(x)$ for small x .

Naive: $1 - \cos(x)$ loses precision. Better:

$$1 - \cos(x) = 2 \sin^2(x/2),$$

which is numerically more stable for small x .

2. Linear algebra essentials

Norm: function $\|\cdot\|$ measuring vector or matrix “size”, satisfying positivity, homogeneity, and the triangle inequality.

Condition number: $\kappa(A) = \|A\| \|A^{-1}\|$; measures sensitivity of the solution of $Ax = b$ to perturbations.

Norms & conditioning.

- Vector p -norm: $\|\mathbf{x}\|_p = (\sum_i |x_i|^p)^{1/p}$.
- Matrix induced norms: $\|A\|_1 = \max_j \sum_i |a_{ij}|$, $\|A\|_\infty = \max_i \sum_j |a_{ij}|$.

Some properties

- Outer product of two non-zero vectors uv^T is always rank 1.
- Product of the eigenvalues is the determinant.
- If columns (or rows) are nearly linearly dependent or have wildly different scales, the matrix is likely ill-conditioned.
- If the Hessian at a critical point has all eigenvalues positive (symmetric case), the Hessian is positive definite and the point is a strict local minimum.

Singular vs nonsingular matrices.

Nonsingular (invertible) $A \in \mathbb{R}^{n \times n}$:

- $\det(A) \neq 0$.
- Inverse exists: $\exists A^{-1}$ with $AA^{-1} = A^{-1}A = I$.
- Columns (and rows) are linearly independent.
- Rank = n ; null space $\mathcal{N}(A) = \{\mathbf{0}\}$.
- For every \mathbf{b} , the system $Ax = \mathbf{b}$ has a unique solution $\mathbf{x} = A^{-1}\mathbf{b}$.
- 0 is not an eigenvalue of A .

Singular $A \in \mathbb{R}^{n \times n}$:

- $\det(A) = 0$.
- No inverse exists.
- Columns (or rows) linearly dependent; rank $< n$.
- Nontrivial null space: $\exists \mathbf{x} \neq \mathbf{0}$ with $A\mathbf{x} = \mathbf{0}$.
- For some \mathbf{b} , $Ax = \mathbf{b}$ has no solution; for others, infinitely many.
- 0 is an eigenvalue: $\exists \mathbf{v} \neq \mathbf{0}$ with $A\mathbf{v} = \mathbf{0}$.

Example (check singularity and implications):

$$A = \begin{bmatrix} 1 & 2 \\ 2 & 4 \end{bmatrix}, \quad \det(A) = 1 \cdot 4 - 2 \cdot 2 = 0 \Rightarrow A \text{ is singular.}$$

So A^{-1} does not exist, columns are multiples, and there are nonzero solutions to $A\mathbf{x} = \mathbf{0}$.

Special matrices (examples and properties).

Identity matrix:

$$I_3 = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}, \quad AI_3 = I_3A = A.$$

Eigenvalues are all 1; $\det(I_3) = 1$.

Diagonal matrix:

$$D = \begin{bmatrix} d_1 & 0 & 0 \\ 0 & d_2 & 0 \\ 0 & 0 & d_3 \end{bmatrix}.$$

- $\det(D) = d_1 d_2 d_3$.

- Eigenvalues are $\lambda_1 = d_1$, $\lambda_2 = d_2$, $\lambda_3 = d_3$.
- $D^{-1} = \text{diag}(1/d_1, 1/d_2, 1/d_3)$ if all $d_i \neq 0$.

Upper triangular matrix:

$$U = \begin{bmatrix} u_{11} & u_{12} & u_{13} \\ 0 & u_{22} & u_{23} \\ 0 & 0 & u_{33} \end{bmatrix}.$$

- $\det(U) = u_{11}u_{22}u_{33}$.
- Eigenvalues are the diagonal entries: $\lambda_i = u_{ii}$.
- Solve $U\mathbf{x} = \mathbf{b}$ by back-substitution in $O(n^2)$.

Lower triangular matrix:

$$L = \begin{bmatrix} \ell_{11} & 0 & 0 \\ \ell_{21} & \ell_{22} & 0 \\ \ell_{31} & \ell_{32} & \ell_{33} \end{bmatrix}.$$

- $\det(L) = \ell_{11}\ell_{22}\ell_{33}$.
- Eigenvalues are ℓ_{ii} .
- Solve $L\mathbf{x} = \mathbf{b}$ by forward substitution.

Symmetric matrix: $A = A^T$; eigenvalues are real and eigenvectors can be chosen orthonormal.

Four fundamental subspaces and solvability.

- Column space $\mathcal{R}(A)$, null space $\mathcal{N}(A)$, row space $\mathcal{R}(A^T)$, left null space $\mathcal{N}(A^T)$.
- Existence: $Ax = \mathbf{b}$ solvable iff $\mathbf{b} \in \mathcal{R}(A)$.
- Uniqueness: solution unique iff $\mathcal{N}(A) = \{\mathbf{0}\}$.

Eigenvalues, eigenvectors, and algorithm sketch.

Eigenpair: (λ, \mathbf{v}) with $\mathbf{v} \neq \mathbf{0}$ such that $A\mathbf{v} = \lambda\mathbf{v}$.

Basic properties:

- $\det(A - \lambda I) = 0$ gives the **characteristic polynomial**; its roots are eigenvalues.
- For each eigenvalue λ , eigenvectors solve $(A - \lambda I)\mathbf{v} = \mathbf{0}$.
- If A has a basis of eigenvectors, then $A = V\Lambda V^{-1}$ with $V = [\mathbf{v}_1 \dots \mathbf{v}_n]$, $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_n)$.

Power method (sketch) for dominant eigenvalue:

1. Choose nonzero \mathbf{x}_0 .
2. For $k = 0, 1, \dots$: compute $\mathbf{y}_{k+1} = A\mathbf{x}_k$, then normalize $\mathbf{x}_{k+1} = \mathbf{y}_{k+1}/\|\mathbf{y}_{k+1}\|$.
3. Rayleigh quotient $\lambda_k = \mathbf{x}_k^T A \mathbf{x}_k$ approximates largest $|\lambda|$.

Practical eigenvalue algorithms: QR iteration with shifts (used by `eig`); repeated orthogonal similarity transforms drive A toward upper triangular (Schur form) whose diagonal entries approximate eigenvalues.

Example (eigenvalues and eigenvectors of a 2×2 matrix).

Let

$$A = \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix}.$$

1) Eigenvalues: solve $\det(A - \lambda I) = 0$.

$$A - \lambda I = \begin{bmatrix} 2 - \lambda & 1 \\ 1 & 2 - \lambda \end{bmatrix}, \quad \det(A - \lambda I) = (2 - \lambda)^2 - 1.$$

Set this to zero:

$$(2 - \lambda)^2 - 1 = 0 \Rightarrow 2 - \lambda = \pm 1 \Rightarrow \lambda_1 = 3, \lambda_2 = 1.$$

2) Eigenvector for $\lambda_1 = 3$: solve $(A - 3I)\mathbf{v}_1 = \mathbf{0}$.

We have

$$A - 3I = \begin{bmatrix} -1 & 1 \\ 1 & -1 \end{bmatrix},$$

so the equation $(A - 3I)\mathbf{v}_1 = \mathbf{0}$ is

$$-v_{1,1} + v_{1,2} = 0 \Rightarrow v_{1,1} = v_{1,2}.$$

Choose $v_{1,1} = 1$, so

$$\mathbf{v}_1 = \begin{bmatrix} 1 \\ 1 \end{bmatrix},$$

and the normalized eigenvector is

$$\hat{\mathbf{v}}_1 = \frac{1}{\sqrt{2}} \begin{bmatrix} 1 \\ 1 \end{bmatrix}.$$

3) Eigenvector for $\lambda_2 = 1$: solve $(A - I)\mathbf{v}_2 = \mathbf{0}$.

We have

$$A - I = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix},$$

so

$$v_{2,1} + v_{2,2} = 0 \Rightarrow v_{2,2} = -v_{2,1}.$$

Choose $v_{2,1} = 1$, so

$$\mathbf{v}_2 = \begin{bmatrix} 1 \\ -1 \end{bmatrix},$$

and the normalized eigenvector is

$$\hat{\mathbf{v}}_2 = \frac{1}{\sqrt{2}} \begin{bmatrix} 1 \\ -1 \end{bmatrix}.$$

Check notation:

$$A\hat{\mathbf{v}}_1 = \lambda_1\hat{\mathbf{v}}_1, \quad A\hat{\mathbf{v}}_2 = \lambda_2\hat{\mathbf{v}}_2.$$

With $V = [\hat{\mathbf{v}}_1 \ \hat{\mathbf{v}}_2]$ and $\Lambda = \text{diag}(\lambda_1, \lambda_2)$, we have

$$A = V\Lambda V^{-1},$$

and since A is symmetric, V can be chosen orthogonal so that $V^{-1} = V^T$.

Orthogonal matrices and projectors

Orthogonal matrix. $Q \in \mathbb{R}^{n \times n}$ is orthogonal if

$$Q^T Q = Q Q^T = I.$$

Properties:

$$Q^{-1} = Q^T, \quad \|Qx\|_2 = \|x\|_2, \quad (Qx) \cdot (Qy) = x \cdot y, \\ \det(Q) = \pm 1, \quad |\lambda_i(Q)| = 1 \text{ for all eigenvalues } \lambda_i.$$

Orthogonal projector. If $Q \in \mathbb{R}^{n \times k}$ has orthonormal columns, then

$$P = Q Q^T$$

is the projector onto $\mathcal{R}(Q)$, with

$$P^T = P, \quad P^2 = P, \quad \lambda_i(P) \in \{0, 1\}, \quad \text{rank}(P) = k.$$

For least squares, with design matrix X ,

$$\hat{y} = X\hat{\theta} = X(X^T X)^{-1} X^T y = P_X y.$$

Positive (semi)definite symmetric matrices

$A \in \mathbb{R}^{n \times n}$ symmetric.

Positive definite (PD):

$$x^T A x > 0 \quad \forall x \neq 0.$$

Positive semidefinite (PSD):

$$x^T A x \geq 0 \quad \forall x.$$

Equivalent facts for symmetric A :

$$A \succ 0 \iff \lambda_i(A) > 0 \forall i, \quad A \succeq 0 \iff \lambda_i(A) \geq 0 \forall i.$$

Cholesky: $A \succ 0 \iff \exists R \text{ upper triangular with } A = R^T R$.

Optimization link. If $H(x)$ (Hessian) is PD at a critical point x^* :

$$\nabla f(x^*) = 0, \quad H(x^*) \succ 0 \Rightarrow x^* \text{ is a strict local minimum (and loc.}$$

Symmetric matrices and eigendecomposition

For symmetric $A \in \mathbb{R}^{n \times n}$, there exists an orthogonal W and real diagonal Λ such that

$$A = W\Lambda W^T, \quad W^T W = I, \quad \Lambda = \text{diag}(\lambda_1, \dots, \lambda_n), \quad \lambda_i \in \mathbb{R}.$$

Useful identities:

$$\det(A) = \prod_{i=1}^n \lambda_i, \quad \text{tr}(A) = \sum_{i=1}^n \lambda_i.$$

If an eigenvalue λ has algebraic multiplicity m and A is symmetric, then its geometric multiplicity is also m , so

$$\dim \mathcal{N}(A - \lambda I) = m, \quad \text{rank}(A - \lambda I) = n - m.$$

Rank-nullity and condition number

For $A \in \mathbb{R}^{m \times n}$,

$$\dim \mathcal{R}(A) + \dim \mathcal{N}(A) = n.$$

If A is invertible, the (matrix) condition number in the 2-norm is

$$\kappa_2(A) = \|A\|_2 \|A^{-1}\|_2 = \frac{\max_i |\lambda_i(A)|}{\min_i |\lambda_i(A)|}.$$

3. SNE and optimization

Jacobian: matrix of first partial derivatives: $J_{ij} = \partial f_i / \partial x_j$.

Newton's method (SNE): iteratively solve linearized system to update guess.

Solve $\mathbf{f}(\mathbf{x}) = \mathbf{0}$:

$$J(\mathbf{x}_k) \Delta \mathbf{x}_k = -\mathbf{f}(\mathbf{x}_k), \quad \mathbf{x}_{k+1} = \mathbf{x}_k + \Delta \mathbf{x}_k.$$

Example (scalar Newton): Solve $x^2 - 2 = 0$.

$$f(x) = x^2 - 2, \quad f'(x) = 2x,$$

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \frac{f(\mathbf{x}_k)}{f'(\mathbf{x}_k)} = \mathbf{x}_k - \frac{x_k^2 - 2}{2x_k}.$$

Gradient: vector of first partial derivatives; direction of steepest ascent.

Hessian: matrix of second partial derivatives; local curvature.

Gradient/Hessian examples.

Quadratic:

$$f(\mathbf{x}) = \frac{1}{2} \mathbf{x}^T Q \mathbf{x} + \mathbf{c}^T \mathbf{x} + d, \quad Q = Q^T,$$

$$\nabla f(\mathbf{x}) = Q\mathbf{x} + \mathbf{c}, \quad H(\mathbf{x}) = Q.$$

Least squares $f(\mathbf{x}) = \frac{1}{2} \|A\mathbf{x} - \mathbf{b}\|^2$:

$$\nabla f(\mathbf{x}) = A^T (A\mathbf{x} - \mathbf{b}), \quad H(\mathbf{x}) = A^T A.$$

Example (minimize quadratic): Solve $\min f(\mathbf{x}) = \frac{1}{2} \mathbf{x}^T Q \mathbf{x} + \mathbf{c}^T \mathbf{x}$. Set gradient $Q\mathbf{x} + \mathbf{c} = 0 \Rightarrow \mathbf{x}^* = -Q^{-1}\mathbf{c}$, assuming $Q \succ 0$.

Lagrangian: $\mathcal{L}(\mathbf{x}, \lambda) = f(\mathbf{x}) + \lambda^T c(\mathbf{x})$ for equality constraints $c(\mathbf{x}) = 0$.

Equality-constrained optimum:

$$\nabla_x \mathcal{L} = \nabla f + J_c(\mathbf{x})^T \lambda = 0, \quad c(\mathbf{x}) = 0.$$

Convex set: $C \subset \mathbb{R}^n$ is convex if

$$x_1, x_2 \in C, t \in [0, 1] \Rightarrow tx_1 + (1-t)x_2 \in C.$$

Convex function: $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is convex if

$$f(tx_1 + (1-t)x_2) \leq tf(x_1) + (1-t)f(x_2), \quad \forall x_1, x_2, t \in [0, 1].$$

If f is twice differentiable and $\nabla^2 f(x) \succeq 0$ for all x , then f is convex. If $\nabla^2 f(x) \succ 0$ for all x , f is *strictly convex* and has at most one minimizer.

Unconstrained first-/second-order conditions: For $\min_x f(x)$ with f differentiable:

$$\nabla f(x^*) = 0 \quad (\text{first-order necessary}).$$

If additionally f is twice differentiable,

$$\nabla f(x^*) = 0, \quad \nabla^2 f(x^*) \succ 0$$

is a *sufficient* condition for x^* to be a strict local minimizer (unique locally).

For convex f , any x^* with $\nabla f(x^*) = 0$ is a *global* minimizer.

Equality-constrained KKT (first order):

$$\min_x f(x) \quad \text{s.t.} \quad c(x) = 0, \quad c : \mathbb{R}^n \rightarrow \mathbb{R}^m.$$

Lagrangian

$$\mathcal{L}(\mathbf{x}, \lambda) = f(\mathbf{x}) + \lambda^T c(\mathbf{x})^T.$$

KKT system (necessary):

$$\nabla_x \mathcal{L}(x^*, \lambda^*) = \nabla f(x^*) + J_c(x^*)^T \lambda^* = 0, \quad c(x^*) = 0.$$

KKT Newton system (matrix form): At iterate (x_k, λ_k) , one Newton step for the KKT conditions solves

$$\begin{bmatrix} \nabla_{xx}^2 \mathcal{L}(x_k, \lambda_k) & J_c(x_k)^T \\ J_c(x_k) & 0 \end{bmatrix} \begin{bmatrix} \Delta x_k \\ \Delta \lambda_k \end{bmatrix} = - \begin{bmatrix} \nabla f(x_k) + J_c(x_k)^T \lambda_k \\ c(x_k) \end{bmatrix},$$

then updates $x_{k+1} = x_k + \Delta x_k$, $\lambda_{k+1} = \lambda_k + \Delta \lambda_k$.

Steepest descent (gradient descent): For unconstrained $\min f(x)$,

$$d_k = -\nabla f(x_k), \quad x_{k+1} = x_k + \alpha_k d_k.$$

- Fixed small $\alpha_k = \alpha \Rightarrow$ guaranteed descent for sufficiently small α , but slow convergence near optimum.
- Exact/line-search α_k (1D minimization along d_k) \Rightarrow better convergence but more work per step.

Newton's method (optimization):

$$d_k = -\nabla^2 f(x_k)^{-1} \nabla f(x_k), \quad x_{k+1} = x_k + d_k.$$

- Quadratic convergence near a solution if $\nabla^2 f(x^*)$ is nonsingular (typically PD). - May diverge or go uphill if $\nabla^2 f$ is indefinite (nonconvex problems).

Quasi-Newton / BFGS: Maintain an approximation $H_k \approx \nabla^2 f(x_k)^{-1}$ updated from gradients only.

$$s_k = x_{k+1} - x_k, \quad y_k = \nabla f(x_{k+1}) - \nabla f(x_k),$$

$$H_{k+1} = \left(I - \frac{s_k y_k^T}{y_k^T s_k} \right) H_k \left(I - \frac{y_k s_k^T}{y_k^T s_k} \right) + \frac{s_k s_k^T}{y_k^T s_k}.$$

Search direction:

$$d_k = -H_k \nabla f(x_k).$$

- Start with $H_0 = I$ (first step = steepest descent). - If $y_k^T s_k > 0$, BFGS preserves positive definiteness of H_k and typically gives superlinear convergence.

Interior-point (barrier) method for inequality constraints: For

$$\min_x f(x) \quad \text{s.t.} \quad h_i(x) \leq 0, \quad i = 1, \dots, m,$$

introduce barrier parameter $\mu > 0$ and barrier function

$$\phi_\mu(x) = f(x) - \mu \sum_{i=1}^m \log(h_i(x)),$$

defined only for strictly feasible x with $h_i(x) < 0$. Algorithm idea:

- For large μ , approximately minimize $\phi_\mu(x)$ (e.g., with Newton).
- Decrease $\mu \rightarrow 0$ and warm-start each new solve from the previous minimizer.

As $\mu \rightarrow 0$, solutions of the barrier problems approach a KKT point of the original constrained problem.

4. Probability basics

Random variable (RV): numerical outcome of a random experiment.

Discrete vs continuous RVs:

- Discrete: X takes values in a countable set $\{x_k\}$, $P(X = x_k) = p_X(x_k)$, $\sum_k p_X(x_k) = 1$.
- Continuous: X has density $f_X(x) \geq 0$, $P(a \leq X \leq b) = \int_a^b f_X(x) dx$, $\int_{-\infty}^{\infty} f_X(x) dx = 1$.

PMF, PDF, CDF (definitions + intuition):

- **PMF (probability mass function)** for discrete X :

$$p_X(k) = P(X = k).$$

Properties: $p_X(k) \geq 0$ for all k ; $\sum_k p_X(k) = 1$. Intuition: a *bar chart of exact probabilities* at each possible value.

- **PDF (probability density function)** for continuous X :

$$f_X(x) \geq 0, \quad \int_{-\infty}^{\infty} f_X(x) dx = 1.$$

Probability over an interval: $P(a \leq X \leq b) = \int_a^b f_X(x) dx$. Properties: nonnegative; integrates to 1; the value $f_X(x)$ is a *density*, not a probability, so $P(X = x) = 0$. Intuition: a *smooth curve of probability intensity*; area under the curve over a region is probability.

- **CDF (cumulative distribution function)** for any X :

$$F_X(x) = P(X \leq x).$$

Discrete: $F_X(x) = \sum_{x_k \leq x} p_X(x_k)$. Continuous: $F_X(x) = \int_{-\infty}^x f_X(t) dt$. Key properties: $0 \leq F_X(x) \leq 1$; nondecreasing in x ; $\lim_{x \rightarrow -\infty} F_X(x) = 0$, $\lim_{x \rightarrow +\infty} F_X(x) = 1$; right-continuous. Intuition: an *accumulated probability curve*, showing how much probability mass has been accumulated up to threshold x .

Expectation: mean; $\mathbb{E}[X]$.

Key formulas.

- $\mathbb{E}[g(X)] = \sum_k g(k)p_X(k)$ or $\int g(x)f_X(x) dx$.
- $\text{Var}(X) = \mathbb{E}[X^2] - (\mathbb{E}[X])^2$.
- $P(A|B) = P(A \cap B)/P(B)$.

- Bayes: $P(A|B) = \frac{P(B|A)P(A)}{P(B)}$.

Marginal and conditional: Joint pmf/pdf $p_{X,Y}(x,y)$:

$$p_X(x) = \sum_y p_{X,Y}(x,y) \quad \text{or} \quad p_X(x) = \int p_{X,Y}(x,y) dy,$$

$$p_{Y|X}(y|x) = \frac{p_{X,Y}(x,y)}{p_X(x)}.$$

Convolution (sum of independent RVs): If $Z = X + Y$ and X, Y independent continuous,

$$f_Z(z) = \int_{-\infty}^{\infty} f_X(x)f_Y(z-x) dx.$$

For discrete,

$$p_Z(k) = \sum_i p_X(i)p_Y(k-i).$$

Covariance and covariance matrix:

$$\text{Cov}(X, Y) = \mathbb{E}[(X - \mathbb{E}X)(Y - \mathbb{E}Y)].$$

For random vector $x \in \mathbb{R}^n$,

$$C = \mathbb{E}[(x - \mu)(x - \mu)^T], \quad \mu = \mathbb{E}[x].$$

C is symmetric, PSD; diagonal entries are variances, off-diagonals are covariances.

Biased vs unbiased estimator: An estimator $\hat{\theta}$ of parameter θ is

unbiased if $\mathbb{E}[\hat{\theta}] = \theta$, biased otherwise.

Examples:

- Sample mean: $\hat{\mu} = \frac{1}{N} \sum_{i=1}^N X_i$ is unbiased for μ .
- Sample variance: $\hat{\sigma}^2 = \frac{1}{N-1} \sum_i (X_i - \hat{\mu})^2$ is unbiased; with $\frac{1}{N}$ it is biased.

Central Limit Theorem (CLT): If X_i are IID with mean μ and variance $\sigma^2 < \infty$, then for large N ,

$$\bar{X}_N = \frac{1}{N} \sum_{i=1}^N X_i \approx \mathcal{N}\left(\mu, \frac{\sigma^2}{N}\right).$$

Only the *sample mean* becomes approximately normal; the histogram of raw samples approaches the true underlying distribution, not necessarily Gaussian.

Example (law of total prob.): Events A_i partition and event B :

$$P(B) = \sum_i P(B|A_i)P(A_i).$$

Notable distributions

Bernoulli (discrete):

$$P(X = 1) = p, \quad P(X = 0) = 1 - p, \quad X \in \{0, 1\}.$$

Mean $\mathbb{E}X = p$, variance $\text{Var}(X) = p(1-p)$.

Binomial (discrete):

$$P(X = k) = \binom{n}{k} p^k (1-p)^{n-k}, \quad k = 0, 1, \dots, n.$$

Mean np , variance $np(1-p)$.

Geometric (discrete, “# of failures before first success” version):

$$P(X = k) = p(1-p)^k, \quad k = 0, 1, 2, \dots$$

Mean $\frac{1-p}{p}$, variance $\frac{1-p}{p^2}$.

Poisson (discrete):

$$P(K = k) = \frac{\lambda^k e^{-\lambda}}{k!}, \quad k = 0, 1, 2, \dots$$

Mean λ , variance λ .

Uniform (continuous on $[a, b]$):

$$f_X(x) = \begin{cases} \frac{1}{b-a}, & a \leq x \leq b, \\ 0, & \text{otherwise,} \end{cases}$$

Mean $\frac{a+b}{2}$, variance $\frac{(b-a)^2}{12}$.

Normal / Gaussian (continuous):

$$f_X(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right), \quad x \in \mathbb{R}.$$

Mean μ , variance σ^2 .

Multivariate Normal (continuous): For $x \in \mathbb{R}^k$,

$$p_X(x) = \frac{1}{(2\pi)^{k/2} \sqrt{\det C}} \exp\left(-\frac{1}{2}(x-\mu)^T C^{-1}(x-\mu)\right),$$

mean vector μ , covariance C .

Exponential (continuous, rate $\lambda > 0$):

$$f_X(x) = \lambda e^{-\lambda x}, \quad x \geq 0.$$

Mean $1/\lambda$, variance $1/\lambda^2$.

Gamma (continuous, shape $\alpha > 0$, rate $\beta > 0$):

$$f_X(x) = \frac{\beta^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\beta x}, \quad x \geq 0.$$

Mean α/β , variance α/β^2 .

5. Models vs data & design

Linear models and estimation

Design matrix X : rows are input vectors \mathbf{x}_i^T , columns correspond to parameters.

Linear model (scalar output):

$$y_i = \mathbf{x}_i^T \boldsymbol{\theta} + \varepsilon_i, \quad \varepsilon_i \sim \mathcal{N}(0, \sigma^2), \\ \mathbf{y} = \mathbf{X}\boldsymbol{\theta} + \boldsymbol{\varepsilon}.$$

Homoscedastic vs heteroscedastic errors:

- Homoscedastic: $\text{Var}(\varepsilon_i) = \sigma^2$ for all i (same variance).
- Heteroscedastic: $\text{Var}(\varepsilon_i)$ depends on i (or on \mathbf{x}_i).

OLS (ordinary least squares): homoscedastic, uncorrelated errors.

$$\hat{\boldsymbol{\theta}}_{\text{OLS}} = \arg \min_{\boldsymbol{\theta}} \sum_i (y_i - \mathbf{x}_i^T \boldsymbol{\theta})^2 = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}.$$

WLS (weighted least squares): known covariance V of ε .

$$\hat{\boldsymbol{\theta}}_{\text{WLS}} = \arg \min_{\boldsymbol{\theta}} (\mathbf{y} - \mathbf{y}^{\text{model}}(\boldsymbol{\theta}))^T V^{-1} (\mathbf{y} - \mathbf{y}^{\text{model}}(\boldsymbol{\theta})).$$

If errors are Gaussian, OLS/WLS coincide with MLE:

Likelihood and MLE:

$$L(\boldsymbol{\theta}) = p(D | \boldsymbol{\theta}), \quad \hat{\boldsymbol{\theta}}_{\text{MLE}} = \arg \max_{\boldsymbol{\theta}} L(\boldsymbol{\theta}) = \arg \max_{\boldsymbol{\theta}} \log L(\boldsymbol{\theta}).$$

For independent Gaussian errors with covariance V ,

$$\log L(\boldsymbol{\theta}) = -\frac{1}{2}(\mathbf{y} - \mathbf{y}^{\text{model}}(\boldsymbol{\theta}))^T V^{-1} (\mathbf{y} - \mathbf{y}^{\text{model}}(\boldsymbol{\theta})) + \text{const.}$$

⇒ maximizing likelihood ≡ minimizing (weighted) SSE.

Example (OLS fit): Data (x_i, y_i) , linear model $y = ax + b$.

$$X = \begin{bmatrix} 1 & x_1 \\ \vdots & \vdots \\ 1 & x_N \end{bmatrix}, \quad \hat{\boldsymbol{\theta}} = \begin{bmatrix} \hat{b} \\ \hat{a} \end{bmatrix} = (X^T X)^{-1} X^T \mathbf{y}.$$

Maximum Likelihood Estimation (MLE): Given data x_1, \dots, x_N assumed IID from a model with parameter $\boldsymbol{\theta}$, the *likelihood* is

$$L(\boldsymbol{\theta}) = \prod_{i=1}^N p(x_i | \boldsymbol{\theta}),$$

viewed as a function of $\boldsymbol{\theta}$. The MLE chooses the parameter value that makes the observed data most probable:

$$\hat{\boldsymbol{\theta}}_{\text{MLE}} = \arg \max_{\boldsymbol{\theta}} L(\boldsymbol{\theta}) = \arg \max_{\boldsymbol{\theta}} \prod_{i=1}^N p(x_i | \boldsymbol{\theta}).$$

Equivalently, because log is monotone,

$$\hat{\boldsymbol{\theta}}_{\text{MLE}} = \arg \max_{\boldsymbol{\theta}} \sum_{i=1}^N \log p(x_i | \boldsymbol{\theta}),$$

which is usually easier to optimize numerically. Intuition: pick the parameter that *best explains the data* under the assumed probabilistic model.

Bayesian vs frequentist, MAP, and Bayes' theorem

Frequentist view: $\boldsymbol{\theta}$ is fixed but unknown; uncertainty is in data. Use point estimates (OLS/WLS/MLE) and confidence or χ^2 regions.

Bayesian view: $\boldsymbol{\theta}$ is random with prior $p(\boldsymbol{\theta})$; data update beliefs via Bayes' theorem to posterior $p(\boldsymbol{\theta} | D)$; use MAP and credible regions.

Bayes' theorem (parameters):

$$p(\boldsymbol{\theta} | D) = \frac{p(D | \boldsymbol{\theta})p(\boldsymbol{\theta})}{p(D)}, \quad p(D) = \int p(D | \boldsymbol{\theta})p(\boldsymbol{\theta}) d\boldsymbol{\theta}.$$

Posterior \propto likelihood \times prior.

Prior/posterior/MAP:

- Prior $p(\boldsymbol{\theta})$: belief about $\boldsymbol{\theta}$ before data.
- Posterior $p(\boldsymbol{\theta} | D) \propto p(D | \boldsymbol{\theta})p(\boldsymbol{\theta})$.
- MAP estimator:

$$\hat{\boldsymbol{\theta}}_{\text{MAP}} = \arg \max_{\boldsymbol{\theta}} p(\boldsymbol{\theta} | D) = \arg \max_{\boldsymbol{\theta}} [\log L(\boldsymbol{\theta}) + \log p(\boldsymbol{\theta})].$$

Connections:

- OLS/WLS/MLE: minimize SSE or weighted SSE; frequentist.
- MAP: regularized SSE (log prior acts as penalty); Bayesian.

Conjugate prior: A prior family $p(\boldsymbol{\theta})$ is conjugate to a likelihood $p(D | \boldsymbol{\theta})$ if the posterior $p(\boldsymbol{\theta} | D)$ is in the same family. Example: $\lambda \sim \text{Gamma}(\alpha, \beta)$ and $k | \lambda \sim \text{Poisson}(\lambda) \Rightarrow \lambda | k \sim \text{Gamma}(\alpha + k, \beta + 1)$.

Fisher information, covariance, and design

Fisher information:

$$\mathcal{I}(\boldsymbol{\theta}) = -\mathbb{E}[\nabla_{\boldsymbol{\theta}}^2 \log L(\boldsymbol{\theta})].$$

For linear, homoscedastic model $\mathbf{y} = \mathbf{X}\boldsymbol{\theta} + \boldsymbol{\varepsilon}$, $\boldsymbol{\varepsilon} \sim \mathcal{N}(0, \sigma^2 I)$:

$$\widehat{\text{Cov}}(\hat{\boldsymbol{\theta}}) \approx \sigma^2 (X^T X)^{-1}, \quad \mathcal{I}(\boldsymbol{\theta}) = \frac{1}{\sigma^2} X^T X.$$

For general nonlinear models with sensitivity Jacobian $J(\boldsymbol{\theta})$ and error covariance V :

$$\mathcal{I}(\boldsymbol{\theta}) \approx J(\boldsymbol{\theta})^T V^{-1} J(\boldsymbol{\theta}).$$

Experimental design: Choose inputs (rows of X or evaluations of J) to reduce parameter uncertainty by optimizing a scalar function of the information matrix.

Let M denote an information matrix (e.g. $M = X^T X$ or $M = J^T J$ for homoscedastic errors). Common criteria:

- A-optimal: minimize $\text{tr}(M^{-1})$ (average parameter variance).
- D-optimal: maximize $\det(M)$ (minimize volume of confidence ellipsoid).
- E-optimal: minimize largest eigenvalue of M^{-1} (worst-case variance).

χ^2 , region of indifference, and credible regions

χ^2 misfit (Gaussian errors): If measurement covariance is V ,

$$\chi^2(\boldsymbol{\theta}) = (\mathbf{y} - \mathbf{y}^{\text{model}}(\boldsymbol{\theta}))^T V^{-1} (\mathbf{y} - \mathbf{y}^{\text{model}}(\boldsymbol{\theta})).$$

Under correct Gaussian noise, $\chi^2(\boldsymbol{\theta}_{\text{true}})$ approximately follows a χ^2 distribution with

$$\nu = (\# \text{ data}) - (\# \text{ fitted parameters}).$$

Region of indifference (frequentist): For significance level α ,

$$\mathcal{R}_{\alpha} = \{\boldsymbol{\theta} : \chi^2(\boldsymbol{\theta}) \leq \chi^2_{\max}\},$$

where χ^2_{\max} is a chosen quantile of the χ^2 distribution (e.g. $1 - \alpha$). Parameters in \mathcal{R}_{α} give fits consistent with the data at level α .

Credible interval/region (Bayesian): A $(1 - \alpha)$ credible set C satisfies

$$P(\theta \in C \mid D) = 1 - \alpha.$$

If $p(\theta \mid D) \approx \mathcal{N}(\hat{\theta}, V)$, an ellipsoidal credible region is

$$(\theta - \hat{\theta})^T V^{-1} (\theta - \hat{\theta}) \leq c_\alpha,$$

with c_α chosen to contain posterior mass $1 - \alpha$. This quadratic form is analogous to a χ^2 threshold (but interpreted Bayesianly).

6. Monte Carlo & MCMC

Monte Carlo integration

Monte Carlo integration (simple sampling): Approximate an integral by averaging random samples.

$$I = \int_{\Omega} f(\mathbf{x}) d\mathbf{x} = \text{Vol}(\Omega) \mathbb{E}[f(X)], \quad X \sim \text{Unif}(\Omega),$$

$$I \approx \frac{\text{Vol}(\Omega)}{N} \sum_{i=1}^N f(\mathbf{x}_i), \quad \mathbf{x}_i \sim \text{Unif}(\Omega).$$

Error scales as $O(N^{-1/2})$, independent of dimension.

Importance sampling: Rewrite

$$I = \int f(x) dx = \int f(x) \frac{p(x)}{p(x)} dx = \mathbb{E}_p \left[\frac{f(X)}{p(X)} \right],$$

where $p(x) > 0$ wherever $f(x) \neq 0$. Estimate:

$$I \approx \frac{1}{N} \sum_{i=1}^N \frac{f(x_i)}{p(x_i)}, \quad x_i \sim p.$$

Choosing p with similar shape to f reduces variance.

Normalized vs unnormalized densities

Normalized vs unnormalized:

- A normalized pdf $\pi(x)$ satisfies $\int \pi(x) dx = 1$.
- Often only known up to a constant: $\tilde{\pi}(x) \propto \pi(x)$.

In many computations (e.g. Metropolis–Hastings), the normalization constant cancels: ratios $\pi(x')/\pi(x)$ can be computed using $\tilde{\pi}$.

Examples:

- Posterior: $\pi(\theta \mid D) \propto p(D \mid \theta)p(\theta)$, evidence $p(D)$ unknown.
- Target for MCMC: only need $\tilde{\pi}$ for acceptance probability.

Markov chains and MCMC

Markov chain: Sequence X_0, X_1, \dots with

$$P(X_{k+1} \in A \mid X_k = x_k, \dots, X_0 = x_0) = P(X_{k+1} \in A \mid X_k = x_k).$$

Target / stationary distribution: A distribution π is stationary for the chain if

$$X_0 \sim \pi \Rightarrow X_k \sim \pi \quad \forall k.$$

A sufficient condition is *detailed balance*:

$$\pi(x)T(x \rightarrow x') = \pi(x')T(x' \rightarrow x) \quad \forall x, x',$$

where $T(x \rightarrow x')$ is the one-step transition density/mass.

MCMC: Construct a Markov chain with stationary distribution $\pi(x)$ (often known only up to a constant), then approximate expectations

$$\mathbb{E}_\pi[g(X)] = \int g(x) \pi(x) dx \approx \frac{1}{N} \sum_{i=B+1}^{B+N} g(x_i),$$

where $\{x_i\}$ is the chain, and B is burn-in.

Markov property (first-order chain): A discrete-time stochastic process $\{X_0, X_1, \dots, X_n\}$ is a (first-order) Markov chain if for all k :

$$\begin{aligned} P(X_{k+1} = x_{k+1} \mid X_k = x_k, X_{k-1} = x_{k-1}, \dots, X_0 = x_0) &= \dots \\ &\dots = P(X_{k+1} = x_{k+1} \mid X_k = x_k) \end{aligned}$$

Intuition: the future depends on the present, not on the full past (“one-step memory”).

Markov chain joint factorization: For a Markov chain X_0, \dots, X_n ,

$$\begin{aligned} P(X_0 = x_0, \dots, X_n = x_n) &= \dots \\ &\dots = P(X_0 = x_0) \prod_{k=1}^n P(X_k = x_k \mid X_{k-1} = x_{k-1}). \end{aligned}$$

If the chain is time-homogeneous with transition matrix P , then $P(X_k = j \mid X_{k-1} = i) = P_{ij}$ and the joint is fully determined by the initial distribution and P .

Bayesian network (directed Markov factorization): Let $X = (X_1, \dots, X_n)$ be random variables with a directed

acyclic graph (DAG) where $\text{Pa}(i)$ are the parents of node i . If the distribution is Markov w.r.t. this DAG, its joint pdf/pmf factorizes as

$$p(x_1, \dots, x_n) = \prod_{i=1}^n p(x_i \mid x_{\text{Pa}(i)}).$$

Intuition: each node only “looks at” its parents; all other conditional independencies follow from the graph.

Metropolis–Hastings (MH) algorithm

Given current state x_k and target $\pi(x)$ (we can use unnormalized $\tilde{\pi}$):

1. Propose $x' \sim q(\cdot \mid x_k)$ (proposal kernel).

2. Compute acceptance probability

$$\alpha(x', x_k) = \min \left(1, \frac{\tilde{\pi}(x') q(x_k \mid x')}{\tilde{\pi}(x_k) q(x' \mid x_k)} \right).$$

3. Draw $u \sim \text{Unif}(0, 1)$; if $u \leq \alpha$, set $x_{k+1} = x'$, else $x_{k+1} = x_k$.

If q is symmetric, $q(x' \mid x) = q(x \mid x')$, then

$$\alpha(x', x) = \min \left(1, \frac{\tilde{\pi}(x')}{\tilde{\pi}(x)} \right).$$

Using MCMC for expectations: If the chain is ergodic (can reach all relevant states and has unique stationary distribution),

$$\frac{1}{N} \sum_{i=B+1}^{B+N} g(x_i) \rightarrow \mathbb{E}_\pi[g(X)]$$

as $N \rightarrow \infty$.

Practical points

- **Burn-in:** discard initial samples before chain reaches stationarity.
- **Thinning:** (optional) keep every k -th sample to reduce autocorrelation; not strictly necessary if N is large.
- **Ergodicity/irreducibility:** proposal must allow reaching all regions where $\pi(x) > 0$; avoid one-sided proposals that can never go backward.
- **Unnormalized targets:** for posterior $p(\theta | D) \propto p(D | \theta)p(\theta)$, use $\tilde{\pi}(\theta) = p(D | \theta)p(\theta)$ directly in MH; no need to compute evidence $p(D)$.

```

1 function samples = metropolis_hastings(logpdf, x0,
2     nsamples, proposal_std)
3 % METROPOLIS_HASTINGS Basic Metropolis-Hastings
4 % sampler
5 %
6 % Inputs:
7 % logpdf - function handle, logpdf(x) returns log p(x)
8 % up to a constant
9 % x0 - initial state (column vector or scalar)
10 % nsamples - number of samples to draw
11 % proposal_std - scalar std dev for Gaussian random-
12 % walk proposal
13 %
14 % Output:
15 % samples - matrix of samples, size = [length(x0),
16 % nsamples]
17 % Preallocate storage for samples
18 dim = numel(x0);
19 samples = zeros(dim, nsamples);
20
21 % Current state and its log-density
22 x_curr = x0(:); % ensure column vector
23 logp_curr = logpdf(x_curr);
24
25 for k = 1:nsamples
26
27     % --- PROPOSE A NEW STATE (Gaussian random
28     % walk) ---
29     % For multivariate x, use independent N(0,
30     % proposal_std^2) in each dim
31     x_prop = x_curr + proposal_std * randn(dim, 1);
32
33     % --- COMPUTE ACCEPTANCE RATIO ---
34     % For symmetric proposals q(x'|x) = q(x|x'),
35     % Hastings term cancels
36     logp_prop = logpdf(x_prop);
37     % log acceptance ratio: log alpha = log p(
38     % x_prop) - log p(x_curr)
39     log_alpha = logp_prop - logp_curr;
40
41 end

```

```

31 % Draw u ~ Uniform(0,1) and accept if log u <
32 % log_alpha
33 if log(rand) < log_alpha
34     % Accept proposal
35     x_curr = x_prop;
36     logp_curr = logp_prop;
37 end
38
39 % Store current state (whether moved or not)
40 samples(:, k) = x_curr;
41

```

7. ODE IVPs

ODE IVP: find $\mathbf{x}(t)$ such that

$$\dot{\mathbf{x}} = f(t, \mathbf{x}), \quad \mathbf{x}(t_0) = \mathbf{x}_0.$$

Time-stepping methods

Explicit Euler (Forward Euler):

$$\mathbf{x}_{n+1} = \mathbf{x}_n + \Delta t f(t_n, \mathbf{x}_n).$$

Implicit Euler (Backward Euler):

$$\mathbf{x}_{n+1} = \mathbf{x}_n + \Delta t f(t_{n+1}, \mathbf{x}_{n+1}),$$

requires solving a nonlinear system each step (e.g. Newton). Good stability for stiff problems.

Midpoint / RK2 (explicit):

$$\mathbf{k}_1 = f(t_n, \mathbf{x}_n), \quad \mathbf{k}_2 = f\left(t_n + \frac{\Delta t}{2}, \mathbf{x}_n + \frac{\Delta t}{2}\mathbf{k}_1\right),$$

$$\mathbf{x}_{n+1} = \mathbf{x}_n + \Delta t \mathbf{k}_2.$$

Higher order than Euler for similar cost.

Classical RK4 (explicit):

$$\begin{aligned} \mathbf{k}_1 &= f(t_n, \mathbf{x}_n), \\ \mathbf{k}_2 &= f\left(t_n + \frac{\Delta t}{2}, \mathbf{x}_n + \frac{\Delta t}{2}\mathbf{k}_1\right), \\ \mathbf{k}_3 &= f\left(t_n + \frac{\Delta t}{2}, \mathbf{x}_n + \frac{\Delta t}{2}\mathbf{k}_2\right), \\ \mathbf{k}_4 &= f\left(t_n + \Delta t, \mathbf{x}_n + \Delta t \mathbf{k}_3\right), \\ \mathbf{x}_{n+1} &= \mathbf{x}_n + \frac{\Delta t}{6} (\mathbf{k}_1 + 2\mathbf{k}_2 + 2\mathbf{k}_3 + \mathbf{k}_4). \end{aligned}$$

Accuracy and local truncation error

Local truncation error (LTE): error made in one step assuming \mathbf{x}_n is exact.

- Explicit Euler: LTE = $O(\Delta t^2)$, global error = $O(\Delta t)$ (1st order).
- Midpoint / RK2: LTE = $O(\Delta t^3)$, global error = $O(\Delta t^2)$ (2nd order).
- RK4: LTE = $O(\Delta t^5)$, global error = $O(\Delta t^4)$ (4th order).

Stability and stiffness

Test problem: $\dot{x} = \lambda x$, exact solution $x(t) = e^{\lambda t} x_0$.

Stability region: set of $z = \lambda \Delta t$ such that the numerical solution does not grow for $\text{Re}(\lambda) < 0$. For a one-step method $x_{n+1} = R(z)x_n$, stability requires $|R(z)| \leq 1$.

Examples:

- Explicit Euler: $R(z) = 1 + z$, stability if $|1 + z| \leq 1$. For real $\lambda < 0$, this gives $-2 \leq \lambda \Delta t \leq 0$.
- Implicit Euler: $R(z) = 1/(1 - z)$; $|R(z)| < 1$ for all $\text{Re}(z) < 0$ (A-stable). Good for stiff problems.

Stiffness: the system has widely separated time scales (eigenvalues with very different negative real parts). Explicit methods require very small Δt for stability; implicit methods allow much larger steps without instability (though accuracy still limits Δt).

Explicit Euler on decay:

$$\dot{x} = -kx, \quad x(0) = x_0.$$

Euler:

$$x_{n+1} = x_n - k \Delta t x_n = (1 - k \Delta t) x_n.$$

Stable if $|1 - k \Delta t| \leq 1 \Rightarrow 0 \leq k \Delta t \leq 2$.

8. ODE BVPs

BVP (boundary value problem): differential equation with conditions at multiple points (typically interval endpoints). Example:

$$y''(x) = F(x, y, y'), \quad x \in [a, b],$$

$$\alpha_1(y(a), y'(a)) = 0, \quad \alpha_2(y(b), y'(b)) = 0.$$

Shooting method (single shooting)

Idea: turn BVP into an IVP with unknown initial condition(s).

- Introduce unknown initial slope/parameters \mathbf{c} (e.g. $y'(a) = c$).
- Solve IVP for given \mathbf{c} to get $y(x; \mathbf{c})$.
- Evaluate boundary residual at $x = b$: $\mathbf{h}(\mathbf{c})$.
- Use a root finder (Newton or secant) to solve $\mathbf{h}(\mathbf{c}) = 0$.

Works well when IVP is well-behaved; can struggle for highly sensitive/unstable IVPs.

Example (linear BVP):

$$y'' = 0, \quad y(0) = 0, \quad y(1) = 1.$$

Let $y'(0) = c$. Solve IVP: $y(x) = cx$. Enforce $y(1) = 1 \Rightarrow c = 1$, so $y(x) = x$.

Finite difference method (FDM) for BVPs

Example:

$$Ay''(x) = f(x), \quad x \in [a, b],$$

with Dirichlet BCs $y(a) = \alpha$, $y(b) = \beta$.

Mesh: $x_i = a + i\Delta x$, $i = 0, \dots, M + 1$ with $\Delta x = (b - a)/(M + 1)$. Use interior points $i = 1, \dots, M$.

Centered second derivative:

$$y''(x_i) \approx \frac{y_{i+1} - 2y_i + y_{i-1}}{\Delta x^2}.$$

Leads to a tridiagonal linear system $A_h \mathbf{y} = \mathbf{b}$ for unknown interior values.

BC handling:

- Dirichlet: set $y_0 = \alpha$, $y_{M+1} = \beta$ directly.
- Neumann: approximate derivative with one-sided or ghost-point finite differences.

Remarks:

- Shooting leverages IVP solvers (adaptive time stepping) but can suffer from sensitivity.
- Finite difference BVP is more robust and leads to linear/nonlinear algebraic systems.

9. PDEs

Classification

General 2D second-order linear PDE:

$$Au_{xx} + Bu_{xy} + Cu_{yy} + Du_x + Eu_y + Gu = H.$$

Classification via $B^2 - 4AC$:

- Elliptic: $B^2 - 4AC < 0$ (e.g. Laplace, Poisson).
- Parabolic: $B^2 - 4AC = 0$ (e.g. heat equation).
- Hyperbolic: $B^2 - 4AC > 0$ (e.g. wave, advection).

Finite difference for elliptic PDEs

Example (Laplace):

$$u_{xx} + u_{yy} = 0 \quad \text{in a rectangle.}$$

Uniform grid with spacing h in both directions; interior node (i, j) :

$$u_{xx}(x_i, y_j) \approx \frac{u_{i+1,j} - 2u_{i,j} + u_{i-1,j}}{h^2},$$

$$u_{yy}(x_i, y_j) \approx \frac{u_{i,j+1} - 2u_{i,j} + u_{i,j-1}}{h^2}.$$

5-point stencil:

$$u_{i+1,j} + u_{i-1,j} + u_{i,j+1} + u_{i,j-1} - 4u_{i,j} = 0.$$

Leads to a sparse linear system $A\mathbf{u} = \mathbf{b}$.

Method of lines (MoL) for time-dependent PDEs

Method of lines: discretize space (e.g. by finite differences) to obtain a large ODE system in time; then apply ODE IVP solvers.

Example: 1D diffusion

$$u_t = Du_{xx}, \quad x \in [0, L],$$

space discretization:

$$\frac{d}{dt} u_j(t) = D \frac{u_{j+1}(t) - 2u_j(t) + u_{j-1}(t)}{\Delta x^2},$$

which is a stiff ODE system that can be integrated with implicit methods or `ode15s`.

Advection, CFL, and upwinding

Example advection:

$$u_t + au_x = 0, \quad a > 0.$$

CFL number:

$$\text{CFL} = \frac{a\Delta t}{\Delta x}.$$

Roughly, Courant–Friedrichs–Lowy condition requires the numerical domain of dependence to contain the PDE's domain of dependence; gives step restrictions for explicit schemes.

Upwind scheme (first-order): For $a > 0$,

$$u_j^{n+1} = u_j^n - \text{CFL} (u_j^n - u_{j-1}^n).$$

This is stable for $0 \leq \text{CFL} \leq 1$.

Central difference in space + forward Euler in time:

$$u_j^{n+1} = u_j^n - \frac{\text{CFL}}{2} (u_{j+1}^n - u_{j-1}^n)$$

is typically unstable for pure advection, motivating upwinding and/or more advanced schemes.

Example (CFL check): $a = 2$, $\Delta x = 0.1$, want CFL = 1:

$$\text{CFL} = \frac{a\Delta t}{\Delta x} = 1 \Rightarrow \Delta t = 0.05.$$

10. Differential-Algebraic Equations (DAEs)

DAE: A system that mixes differential and algebraic equations.

Semi-explicit form:

$$\dot{\mathbf{x}} = f(t, \mathbf{x}, \mathbf{y}), \quad 0 = g(t, \mathbf{x}, \mathbf{y}),$$

where \mathbf{x} are differential states (time derivatives appear) and \mathbf{y} are algebraic states (no time derivatives).

Implicit form:

$$F(t, \mathbf{z}, \dot{\mathbf{z}}) = 0,$$

where \mathbf{z} collects all states. If $\partial F/\partial \dot{\mathbf{z}}$ is singular, the system is a genuine DAE (not just an implicit ODE).

Mass-matrix form:

$$M(t, \mathbf{y}) \dot{\mathbf{y}} = F(t, \mathbf{y}),$$

with (possibly singular) mass matrix M . If M is non-singular, this is an ODE; if M is singular, it is a DAE suitable for mass-matrix solvers (e.g. `ode15s`).

Differential index: The smallest number of times the algebraic equations must be differentiated (with respect to time) to obtain an explicit ODE for all derivatives $\dot{\mathbf{x}}, \dot{\mathbf{y}}$.

Index-1 test (semi-explicit): Given

$$\dot{\mathbf{x}} = f(t, \mathbf{x}, \mathbf{y}), \quad 0 = g(t, \mathbf{x}, \mathbf{y}),$$

differentiate $g = 0$ once:

$$\frac{d}{dt}g(t, \mathbf{x}, \mathbf{y}) = g_t + g_{\mathbf{x}}\dot{\mathbf{x}} + g_{\mathbf{y}}\dot{\mathbf{y}} = 0.$$

If $g_{\mathbf{y}}$ is nonsingular, solve for $\dot{\mathbf{y}}$:

$$\dot{\mathbf{y}} = -g_{\mathbf{y}}^{-1}(g_t + g_{\mathbf{x}}\dot{\mathbf{x}}),$$

so the system is index-1. If further differentiation is needed to isolate derivatives, the index is > 1 .

Implicit (hidden) constraints: Higher-index DAEs have additional constraints obtained by differentiating $g(t, \mathbf{x}, \mathbf{y}) = 0$ (e.g. $\dot{g} = 0$, $\ddot{g} = 0$). Solutions must satisfy both the explicit algebraic constraints $g = 0$ and these derivative constraints. Index-1 DAEs typically do not have nontrivial hidden constraints.

Consistent initialization: An initial triple $(\mathbf{x}_0, \dot{\mathbf{x}}_0, \mathbf{y}_0)$ that satisfies

- the DAE at t_0 :

$$\dot{\mathbf{x}}_0 = f(t_0, \mathbf{x}_0, \mathbf{y}_0), \quad 0 = g(t_0, \mathbf{x}_0, \mathbf{y}_0),$$

- and, for higher-index systems, any implied derivative constraints obtained by differentiating $g = 0$ enough times (e.g. $\dot{g}(t_0, \mathbf{x}_0, \mathbf{y}_0, \dot{\mathbf{x}}_0, \dot{\mathbf{y}}_0) = 0$).

Not every choice of $(\mathbf{x}_0, \mathbf{y}_0)$ is allowed; algebraic and hidden constraints restrict admissible initial conditions.

Example (semi-explicit index-1):

$$\dot{x} = y, \quad 0 = x + y.$$

Algebraic constraint: $x + y = 0 \Rightarrow y = -x$. Substitute into the differential equation:

$$\dot{x} = -x,$$

which is an ODE. The system is index-1. Consistent initialization requires $y_0 = -x_0$.

Variables, equations, and DoF (semi-explicit index-1): Consider

$$\dot{\mathbf{x}} = f(t, \mathbf{x}, \mathbf{y}), \quad 0 = g(t, \mathbf{x}, \mathbf{y}),$$

with differential states $\mathbf{x} \in \mathbb{R}^{n_d}$, algebraic states $\mathbf{y} \in \mathbb{R}^{n_a}$, and $g_{\mathbf{y}}$ full rank so that g provides n_a independent algebraic equations. At t_0 :

- Unknowns in the initialization: $(\mathbf{x}_0, \mathbf{y}_0) \in \mathbb{R}^{n_d+n_a}$.
- Algebraic constraints: $g(t_0, \mathbf{x}_0, \mathbf{y}_0) = 0$ impose n_a independent conditions.

Thus the number of degrees of freedom in specifying a consistent initialization is

$$\text{DoF} = (n_d + n_a) - n_a = n_d.$$

A convenient viewpoint for regular index-1 systems is:

- choose the n_d components of \mathbf{x}_0 freely,
- then solve $g(t_0, \mathbf{x}_0, \mathbf{y}_0) = 0$ for the n_a components of \mathbf{y}_0 to obtain a consistent $(\mathbf{x}_0, \mathbf{y}_0)$.

For higher-index DAEs, additional hidden constraints from differentiated g further reduce admissible initial values and may require specifying some derivatives at t_0 as well.

Index reduction (idea): Rewrite a higher-index DAE as an equivalent index-1 DAE by:

- differentiating algebraic equations to expose hidden constraints,
- eliminating redundant variables/constraints or introducing new variables for derivatives,
- reformulating the model in semi-explicit or mass-matrix index-1 form.

This facilitates numerical solution with standard DAE solvers.

Connection to ODE-IVPs:

- Index-0 DAE \equiv ODE-IVP: $\dot{\mathbf{x}} = f(t, \mathbf{x})$.
- Index-1 DAEs can typically be solved directly with DAE-capable time-stepping methods once consistent initial conditions are available.
- Higher-index DAEs are usually reduced to index-1 before numerical solution.

Numerical solution in MATLAB (index-1 DAEs)

Mass-matrix solvers (ode15s, ode23t): Solve

$$M(t, \mathbf{y}) \dot{\mathbf{y}} = F(t, \mathbf{y}),$$

with possibly singular M . You provide M , F , an initial \mathbf{y}_0 (satisfying algebraic constraints), and optionally $\dot{\mathbf{y}}_0$; the solver can also attempt to determine a consistent $\dot{\mathbf{y}}_0$.

Fully implicit solver (ode15i): Solve

$$F(t, \mathbf{y}, \dot{\mathbf{y}}) = 0$$

for index-1 DAEs in implicit form. Typical workflow:

- Define a function for $F(t, \mathbf{y}, \dot{\mathbf{y}})$.
- Use `decic` to compute consistent $(\mathbf{y}_0, \dot{\mathbf{y}}_0)$ from approximate guesses.
- Call `ode15i` with F , $(\mathbf{y}_0, \dot{\mathbf{y}}_0)$, and the time grid.

11. Quick pattern-matching table

Question type	Immediate move
Solve $Ax = b$	Check $m, n, \kappa(A)$; use backslash; think about uniqueness.
Nonlinear equations	Form $f(x)$, use Newton-/tangent iteration; need Jacobian and stopping criteria.
Unconstrained min	Write $\nabla f, H$; use Newton-/gradient descent; check $H \succ 0$.
Regression	Build X, y ; use OLS/WLS; compute θ , covariance, goodness-of-fit.
MLE / MAP	Write likelihood/log-likelihood; differentiate, set to 0; add prior for MAP.
MC integral	Express integral as expectation; choose proposal/target; define estimator.
ODE-IVP	Identify stiffness; choose explicit vs implicit; argue stability and accuracy.
BVP	BVP vs IVP; choose shooting or FD/FEM; formulate equations.
PDE (steady)	Classify; discretize in space; build sparse linear system.
PDE (time)	Use method of lines or explicit scheme; compute CFL; choose upwinding if needed.
DAE	Write semi-explicit/mass form; estimate index; enforce consistent init.

12. Useful MATLAB functions

Linear algebra:

- `A\b`: solve $Ax = b$.
- `eig(A)`: eigenvalues/vectors.
- `svd(A)`: SVD, rank, cond.
- `norm(A,p)`, `cond(A)`.

Optimization / SNE:

- `fsolve`: nonlinear systems.
- `fminunc`: unconstrained min.
- `fmincon`: constrained min.

Probability / stats:

- `rand`, `randn`.
- `binornd`, `poissrnd`.
- `normpdf`, `normcdf`.

ODEs / DAEs:

- `ode45`: nonstiff IVPs.
- `ode23s`, `ode15s`: stiff IVPs.
- `ode15i`: implicit ODE/index-1 DAE.

Sparse / linear systems:

- `sparse`: create sparse matrices.
- `gmres`, `bicgstab`: iterative solvers.

13. Exam patterns & tricks

Linear algebra & discretization

- Always mention conditioning: ill-conditioned matrices amplify data/round-off error; well-conditioned ones do not, regardless of solver. Tridiagonal/banded systems from 1D FDM/FVM solve in $\mathcal{O}(N)$ vs. $\mathcal{O}(N^3)$ for dense.
- For “does a solution exist / is it unique?”, map to four subspaces: existence $\Leftrightarrow b \in \text{range}(A)$; uniqueness $\Leftrightarrow \text{null}(A) = \{0\}$ (full column rank).

ODE IVPs

- For stability questions, immediately apply the test problem $\dot{x} = \lambda x$ and write $x_{n+1} = R(z)x_n$, $z = \lambda\Delta t$. Stability region is $\{z : |R(z)| \leq 1\}$. Forward Euler: $R(z) = 1 + z$; Backward Euler: $R(z) = 1/(1 - z)$.
- For accuracy, say: explicit/implicit Euler are first-order in Δt ; midpoint / trapezoidal are second-order; RK4 is fourth-order. Global error order is one less than the local truncation error order.
- Stiffness: if eigenvalues span many orders of magnitude in negative real part, explicit methods need $\Delta t \ll 1/|\lambda_{\max}|$ purely for stability; implicit methods permit much larger Δt for the same problem.

ODE BVPs

- Shooting method: define the scalar residual $h(c)$ (e.g. $h(c) = y(b; c) - y_b$), then view it as a 1D root-finding problem. Newton in c typically needs $h'(c)$ via sensitivities or finite differences.
- Finite differences: interior points use centered stencils giving $\mathcal{O}(\Delta x^2)$ truncation error; if your Taylor expansion yields $\mathcal{O}(\Delta x)$ you almost certainly misapplied the stencil or indices.
- Think of linear BVPs as sparse linear systems: interior nodes give a tridiagonal (1D) or banded (2D) matrix; mention using backslash on sparse matrices for $\mathcal{O}(N)$ - $\mathcal{O}(N \text{ band}^2)$ cost.

PDEs

- Classification: look only at second-derivative coefficients A, B, C in $Au_{xx} + Bu_{xy} + Cu_{yy} + \dots = 0$. If $B^2 - 4AC < 0$ elliptic, $= 0$ parabolic, > 0 hyperbolic. Lower-order terms do not affect type.
- FDM/FVM: “balance form” answers are valued. Always write: accumulation = in-flux – out-flux + reaction; then insert 2-point fluxes. This is often enough for most of the credit even if algebra gets messy.
- Method of lines: say “discretize space only, get large ODE-IVP system in time, then choose explicit/implicit ODE solver based on stiffness (e.g. `ode45` vs. `ode15s`).” Mention this pipeline explicitly.
- Von Neumann / CFL: plug Fourier mode e^{ikx} into the scheme to get growth factor $G(k)$; require $|G(k)| \leq 1$ for all k . For upwind advection, recover $0 \leq \text{CFL} \leq 1$; central advection + FE is typically unstable.

DAEs

- Index trick: for semi-explicit DAEs $x' = f(t, x, y)$, $0 = g(t, x, y)$, differentiate algebraic equations until you can solve for all derivatives; the number of differentiations is the index (ODE-IVPs are index-0).
- Consistent initialization: check both explicit algebraic constraints and any hidden ones revealed by differentiation at t_0 . If (x_0, y_0) does not satisfy all, the ICs are inconsistent.
- Mass-matrix form: if $M\dot{z} = F(t, z)$ with singular M , say “this is a DAE; use an implicit stiff solver (e.g. `ode15s` / `ode15i`) and supply a consistent initial condition, possibly computed with a ‘decic’-style routine.”

Probability, estimation, MCMC

- Likelihood vs. prior vs. posterior: exams love “write the likelihood” and “derive MLE / MAP” for simple models (normal, Poisson, exponential). Always write $\text{posterior} \propto \text{likelihood} \times \text{prior}$ and normalize only if asked.
- For linear regression / experimental design, tie “better experiment” answers to the design matrix X : more spread in inputs and less collinearity \Rightarrow more informative (smaller parameter covariance).
- MCMC “integral form”: the standard answer is $\frac{\int g(x)f(x)dx}{\int f(x)dx}$. Recognize code patterns where samples are used to estimate expectations of this form; mention unbiasedness as $M \rightarrow \infty$.

14. Taylor Expansion

A Taylor expansion of a sufficiently smooth scalar function f about a point x_0 is

$$f(x) \approx \sum_{k=0}^n \frac{f^{(k)}(x_0)}{k!} (x - x_0)^k.$$

For specific orders:

$$\text{Order 1: } f(x) \approx f(x_0) + f'(x_0)(x - x_0),$$

$$\text{Order 2: } f(x) \approx f(x_0) + f'(x_0)(x - x_0) + \frac{1}{2} f''(x_0)(x - x_0)^2,$$

$$\begin{aligned} \text{Order 3: } f(x) \approx & f(x_0) + f'(x_0)(x - x_0) + \frac{1}{2} f''(x_0)(x - x_0)^2 + \\ & \dots + \frac{1}{6} f^{(3)}(x_0)(x - x_0)^3. \end{aligned}$$

Finite difference coefficients

Central finite differences (uniform h)

Deriv.	Acc.	-5	-4	-3	-2	-1	0	1	2	3	4	5
1	2				-1/2	0	1/2					
1	4				1/12	-2/3	0	2/3	-1/12			
1	6			-1/60	3/20	-3/4	0	3/4	-3/20	1/60		
1	8		1/280	-4/105	1/5	-4/5	0	4/5	-1/5	4/105	-1/280	
2	2				1	-2	1					
2	4				-1/12	4/3	-5/2	4/3	-1/12			
2	6			1/90	-3/20	3/2	-49/18	3/2	-3/20	1/90		
2	8		-1/560	8/315	-1/5	8/5	-205/72	8/5	-1/5	8/315	-1/560	
3	2				-1/2	1	0	-1	1/2			
3	4			1/8	-1	13/8	0	-13/8	1	-1/8		
3	6		-7/240	3/10	-169/120	61/30	0	-61/30	169/120	-3/10	7/240	
4	2				1	-4	6	-4	1			
4	4			-1/6	2	-13/2	28/3	-13/2	2	-1/6		
4	6		7/240	-2/5	169/60	-122/15	91/8	-122/15	169/60	-2/5	7/240	
5	2			-1/2	2	-5/2	0	5/2	-2	1/2		
5	4			1/6	-3/2	13/3	-29/6	0	29/6	-13/3	3/2	-1/6
5	6	-13/288	19/36	-87/32	13/2	-323/48	0	323/48	-13/2	87/32	-19/36	13/288
6	2				1	-6	15	-20	15	-6	1	
6	4			-1/4	3	-13	29	-75/2	29	-13	3	-1/4
6	6	13/240	-19/24	87/16	-39/2	323/8	-1023/20	323/8	-39/2	87/16	-19/24	13/240

Forward finite differences (uniform h)

Deriv.	Acc.	0	1	2	3	4	5	6	7	8
1	1	-1	1							
1	2	-3/2	2	-1/2						
1	3	-11/6	3	-3/2	1/3					
1	4	-25/12	4	-3	4/3	-1/4				
1	5	-137/60	5	-5	10/3	-5/4	1/5			
1	6	-49/20	6	-15/2	20/3	-15/4	6/5	-1/6		
2	1	1	-2	1						
2	2	2	-5	4	-1					
2	3	35/12	-26/3	19/2	-14/3	11/12				
2	4	15/4	-77/6	107/6	-13	61/12	-5/6			
2	5	203/45	-87/5	117/4	-254/9	33/2	-27/5	137/180		
2	6	469/90	-223/10	879/20	-949/18	41	-201/10	1019/180	-7/10	
3	1	-1	3	-3	1					
3	2	-5/2	9	-12	7	-3/2				
3	3	-17/4	71/4	-59/2	49/2	-41/4	7/4			
3	4	-49/8	29	-461/8	62	-307/8	13	-15/8		
3	5	-967/120	638/15	-3929/40	389/3	-2545/24	268/5	-1849/120	29/15	
3	6	-801/80	349/6	-18353/120	2391/10	-1457/6	4891/30	-561/8	527/30	-469/240
4	1	1	-4	6	-4	1				
4	2	3	-14	26	-24	11	-2			
4	3	35/6	-31	137/2	-242/3	107/2	-19	17/6		
4	4	28/3	-111/2	142	-1219/6	176	-185/2	82/3	-7/2	
4	5	1069/80	-1316/15	15289/60	-2144/5	10993/24	-4772/15	2803/20	-536/15	967/240

Backward finite differences (uniform h)

Deriv.	Acc.	-8	-7	-6	-5	-4	-3	-2	-1	0	
1	1							-1	1		
1	2						1/2	-2	3/2		
1	3					-1/3	3/2	-3	11/6		
2	1							1	-2	1	
2	2						-1	4	-5	2	
3	1						-1	3	-3	1	
3	2					3/2	-7	12	-9	5/2	
4	1					1	-4	6	-4	1	
4	2					-2	11	-24	26	-14	3